

AI for Economic Research: Dynamic Models, Language, and Agents

Lecture 4.3: Euler-Equation Method & PyTorch Implementation

● Zhigang Feng

2026-07-06 04:29:46

Topic 4.3 Agenda

1. **Euler Equation Method** — supervised-learning formulation
2. **PyTorch Implementation** — code walkthrough

Arc: *T1 Motivation* → *T2 DL VFI + AC preview* → *T3 Euler + PyTorch* → *T4 Structured NNs + Validation*

Euler Equation Method + Supervised Learning

Context: Two Approaches to the Same Problem

- We have seen the **actor-critic** approach: train two networks to satisfy the Bellman equation
- Now we consider an alternative: **Euler equation minimization**
 - Derive the first-order conditions (Euler equations)
 - Train a *single* policy network to satisfy them
- **Comparison:**
 - Actor-Critic (Bellman): Two networks, no need to derive FOCs, RL flavor
 - Euler Equation: One network, requires deriving FOCs, supervised learning flavor
- Both are valid; which is better depends on the model:
 - Euler equation method is very precise when FOCs are available
 - Actor-Critic is more general (handles kinks, discrete choices)

*Euler-equation training is a **supervised** problem (regress the FOC residual against zero); actor-critic is **reinforcement** (policy gradient on the Bellman objective). Lec05 T3 formalizes the actor-critic loss derivation; here we treat both as PyTorch optimization.*

The Euler Equation: Optimality Condition

- Capital accumulation: $k' = g(k)$, consumption: $c = f(k) + (1 - \delta)k - g(k)$
- For CRRA utility $u(c) = \frac{c^{1-\sigma}}{1-\sigma}$ and Cobb-Douglas production, the equilibrium satisfies:

$$u'(c) = \beta u'(c') [\alpha g(k)^{\alpha-1} + 1 - \delta] \quad (\text{EE})$$

- **Interpretation:** The marginal cost of saving today equals the discounted marginal benefit tomorrow
- **Goal:** Learn a policy g that makes the **Euler residual** vanish:

$$\mathcal{R}(k; g) = u'(c) - \beta u'(c') [\alpha g(k)^{\alpha-1} + 1 - \delta]$$

for all economically relevant k

Euler Equation: Necessary but Not Sufficient

- **Important caveat:** The Euler equation is a *necessary* condition for optimality, but **not sufficient** by itself
- For a complete characterization, we also need:
 - **Transversality condition:** $\lim_{t \rightarrow \infty} \beta^t u'(c_t) k_{t+1} = 0$
 - **Feasibility constraints:** $c \geq 0, k' \geq 0$
- **Handling inequality constraints:**
 - Kuhn-Tucker complementary slackness conditions arise
 - Can be converted to equalities via Fischer-Burmeister functions or penalty methods
 - Or handled via architectural constraints (sigmoid output)
- For now, we focus on **interior solutions** where the Euler equation holds with equality throughout

From Euler Residual to Loss Function

- Approximate the policy by a neural network: $g(k) \approx \Gamma_{\gamma_g}(k)$
- Sample $k_i \sim D(k)$ (e.g., stationary distribution or uniform)
- Define the **mean-squared Euler residual** loss as an expectation over the sampling distribution:

$$\mathcal{L}_g(\gamma_g) = \mathbb{E}_{k \sim D(k)} [\mathcal{R}(k; \Gamma_{\gamma_g})^2] \approx \frac{1}{N} \sum_{i=1}^N [\mathcal{R}(k_i; \Gamma_{\gamma_g})]^2$$

The empirical mean $\frac{1}{N} \sum_i$ is the special case of $\mathbb{E}_{k \sim D(k)}$ using N i.i.d. draws.

- **This is supervised learning:** the “label” for each state k_i is zero (the Euler equation should hold exactly)
- Optimal policy parameters satisfy: $\nabla_{\gamma_g} \mathcal{L}_g(\gamma_g^*) = 0$
- All derivatives computed automatically via autograd

The Mean Value Intuition

- We minimize the *average* Euler residual across sampled states:

$$\mathcal{L}_g(\gamma_g) = \mathbb{E}_{k \sim D(k)} [\mathcal{R}(k; \Gamma_{\gamma_g})^2] \approx \frac{1}{N} \sum_{i=1}^N \mathcal{R}(k_i; \Gamma_{\gamma_g})^2$$

- **Intuition:**

- If a candidate policy satisfies the Euler equation *on average* across all sampled states, it should be a good approximation to the true equilibrium policy
- If the average residual is large, the policy must be violating optimality somewhere
- This “average fitness” criterion:
 - Ensures global accuracy across the domain
 - Is computationally tractable via Monte Carlo sampling
 - Naturally integrates with gradient-based NN training

The Stochastic Case: Expectation in the Euler Equation

- With a productivity shock z , the Euler equation carries a **conditional expectation** over next period's shock z' :

$$u'(c_t) = \beta \mathbb{E}_{z'|z} [u'(c_{t+1})(\alpha k_{t+1}^{\alpha-1} z_{t+1} + 1 - \delta)]$$

- The residual $\mathcal{R}$ now depends on that inner expectation — there is no closed form.
- Monte-Carlo estimate:** replace $\mathbb{E}_{z'|z}$ by an average over J shock draws z'_{ij} :

$$\hat{\mathcal{R}}(k_i; g) = u'(c_i) - \frac{\beta}{J} \sum_{j=1}^J u'(c'_{ij}) [\alpha k_i^{\alpha-1} z'_{ij} + 1 - \delta]$$

- Two layers of sampling:** an *outer* average over states $k_i \sim D(k)$ (the loss $\mathbb{E}_{k \sim D(k)}$) and an *inner* average over shocks z'_{ij} (this conditional expectation).

Variance Reduction: Gauss–Hermite Quadrature

- Plain Monte-Carlo over J random shock draws is *noisy*: its error falls only like $1/\sqrt{J}$, and that noise enters every gradient.
- **Gauss–Hermite quadrature** replaces the random draws with a few *deterministic* nodes z'_j and weights w_j :

$$\mathbb{E}_{z'|z}[h(z')] = \int h(z') \phi(z') dz' \approx \sum_{j=1}^J w_j h(z'_j),$$

exact whenever h is a polynomial of degree $\leq 2J-1$ (against the Gaussian weight ϕ).

- Log-normal shocks $z' = e^{\mu + \sigma \varepsilon}$, $\varepsilon \sim \mathcal{N}(0, 1)$: change variables to the standard Gaussian, then read nodes/weights from a table.
- **Why it works**: smooth integrands are nearly polynomial, so a handful of well-placed nodes ($J \approx 5-7$) match thousands of random draws — same accuracy, *zero* sampling noise.
- **Why the loss stays differentiable**: the nodes and weights are *constants*, so the loss is a fixed weighted sum of smooth network outputs — autograd differentiates straight through. Mini-batching over states only changes *which* states enter the sum; it never breaks the gradient (we sample exogenous shocks/states, not policy-dependent randomness, so no reparameterization trick is needed).

Neural Network for the Policy Function

- Two-hidden-layer MLP (64–64–1) with ReLU and Sigmoid:

$$\Gamma_{\gamma_g}(k) = \sigma_2\left(W_2 \cdot \sigma_1(W_1 k + b_1) + b_2\right)$$

- Hidden activation $\sigma_1 = \text{ReLU}$; output activation $\sigma_2 = \text{Sigmoid}$
- Sigmoid output is scaled to $[0, f(k) + (1 - \delta)k - \varepsilon]$ to ensure feasibility:

$$0 \leq \Gamma_{\gamma_g}(k) \leq f(k) + (1 - \delta)k$$

- All derivatives $\partial \mathcal{L}_g / \partial \gamma_g$ are computed automatically by autograd
- No need to derive or code analytical Jacobians

Feasibility by Construction

- **Why Sigmoid, then rescale?**

- Sigmoid output $s \in (0, 1)$ ensures $g(k) = s \cdot [f(k) + (1 - \delta)k - \varepsilon]$ stays in the feasible interior
- Avoids generating samples with $c \leq 0$ during training (which would blow up $u'(c) = c^{-\sigma}$)

- **Why subtract ε ?** A small slack prevents marginal-utility divergence near $c = 0$

- **Alternatives:**

- Softplus output: always positive, smooth
- Hard `torch.clamp`: simple but kills gradients at the boundary
- Log-parameterization: train on $\log c$ rather than c

- Architectural feasibility outperforms post-hoc clipping — Lab 3B verifies this empirically

Minimising the Euler Residual: Training Loop

- Initialize $\gamma_g^{(0)}$; choose learning rate α_g
- **Repeat until convergence:**
 1. Draw a mini-batch $\{k_i\}_{i=1}^B$
 2. Compute loss $\mathcal{L}_g(\gamma_g)$ and its gradient via back-propagation
 3. Gradient step:

$$\gamma_g^{(j+1)} = \gamma_g^{(j)} - \alpha_g \nabla_{\gamma_g} \mathcal{L}_g(\gamma_g^{(j)})$$

4. Optionally resample k_i from the same distribution $D(k)$
- **Pros of Euler method:**
 - Very precise policy functions; theoretically grounded
 - Only one network to train (simpler than actor-critic)
 - **Cons:** Does not yield the value function V ; requires deriving FOCs

PyTorch Implementation

© Zhigang Feng

Redistribution prohibited without permission

Implementation Overview: Two Approaches

- **Approach 1 — Deep VFI (Lab 3B, Part 1):** PolicyNet + ValueNet, T -step lookahead with soft constraints; result = $g(k)$ and $V(k)$.
- **Approach 2 — Euler Equation (Lab 3B, Part 2):** Single PolicyNetwork; loss = mean-squared Euler residual; result = policy $g(k)$.
- Both validated against the analytical solution from Lab 3A.

Common PyTorch skeleton (both approaches)

1. **Define networks:** PolicyNet (Sigmoid output, feasibility-rescaled) and/or ValueNet.
2. **Build loss:** Bellman MSE = $(V(k) - (u + \beta V(k')))^2$ and/or Euler residual = $u'(c) - \beta \mathbb{E}[u'(c') (\alpha k'^{\alpha-1} z' + 1 - \delta)]$.
3. **Train:** AdamW on mini-batches sampled from $D(k)$; alternate Critic/Actor steps when both nets are present.

Network Architecture: PolicyNet and ValueNet

Code targets *PyTorch* ≥ 2.3 (2024); AdamW optimizer per Lec03 T3:225. Lab 3A/3B notebooks use matching versions.

```
1 class ValueNet(nn.Module):
2     def __init__(self, nh=64):
3         super().__init__()
4         self.net = nn.Sequential(
5             nn.Linear(1, nh), nn.ReLU(),
6             nn.Linear(nh, nh), nn.ReLU(),
7             nn.Linear(nh, 1))
8     def forward(self, k): return self.net(k)
9
10 class PolicyNet(nn.Module):
11     def __init__(self, alpha, eps, nh=64):
12         super().__init__()
13         self.alpha, self.eps = alpha, eps
14         self.net = nn.Sequential(
15             nn.Linear(1, nh), nn.ReLU(),
16             nn.Linear(nh, nh), nn.ReLU(),
17             nn.Linear(nh, 1), nn.Sigmoid()) # Output in [0, 1]
18     def forward(self, k):
19         max_kp = torch.clamp(k**self.alpha - self.eps, 1e-9)
20         return self.net(k) * max_kp # Scale to [0, f(k)-eps]
```

- Sigmoid ensures $g(k) \in [0, f(k) - \varepsilon]$: feasibility by construction
- ValueNet: unconstrained output (value can be any real number)

Policy Evaluation: Bellman MSE

```
1 def value_update_step(self):
2     self.value_optimizer.zero_grad()
3     # Sample batch of capital states
4     k_batch = torch.rand(self.n_batch, 1, device=self.device) \
5         * (self.k_max - self.k_min) + self.k_min
6
7     v_current = self.value_net(k_batch)           # V(k)
8     kp_batch = self.policy_net(k_batch)           # k' = g(k)
9     u_batch = self.utility(k_batch, kp_batch)     # u(c)
10    v_next = self.value_net(kp_batch).detach()     # V(k') -- detach!
11    bellman_rhs = u_batch + self.beta * v_next     # u + beta*V(k')
12
13    loss = nn.functional.mse_loss(v_current, bellman_rhs)
14    loss.backward()
15    self.value_optimizer.step()
16    return loss.item()
```

- Minimizes $\mathbb{E}[(V(k) - (u + \beta V(k')))^2]$
- `detach()`: prevents gradients from flowing through $V(k')$ — treats the target as fixed

Policy Improvement: Gradient Ascent

```
1 def policy_step(self):
2     self.policy_optimizer.zero_grad()
3     # Sample batch of capital states
4     k_batch = torch.rand(self.n_batch, 1, device=self.device) \
5         * (self.k_max - self.k_min) + self.k_min
6
7     kp_batch = self.policy_net(k_batch)          # k' = g(k)
8     u_batch = self.utility(k_batch, kp_batch)    # u(c)
9     v_next = self.value_net(kp_batch)           # V(k')
10
11     objective = torch.mean(u_batch + self.beta * v_next)
12
13     loss = -objective    # Maximize objective = Minimize -objective
14     loss.backward()
15     self.policy_optimizer.step()
16     return objective.item()
```

- Maximizes $\mathbb{E}[u(k, g(k)) + \beta V(g(k))]$
- Gradients flow through **both** u and V into the policy network parameters

The Training Loop: Alternating Updates

```
1 def train(self):
2     print(f"Training started on device={self.device}")
3     pbar = tqdm(range(self.n_epoch), desc="Outer Loop")
4
5     for _ in pbar:
6         # Phase 1: Update Value Net (several inner steps)
7         for _ in range(self.n_inner_value):
8             v_loss = self.value_update_step()
9             self.value_losses.append(v_loss)
10
11        # Phase 2: Update Policy Net (several inner steps)
12        for _ in range(self.n_inner_policy):
13            policy_obj = self.policy_update_step()
14            self.policy_objectives.append(policy_obj)
15
16        pbar.set_postfix({
17            "ValueLoss": f"{v_loss:.4e}",
18            "PolicyObj": f"{policy_obj:.4e}"
19        })
```

- Alternates between updating the Critic and the Actor
- Multiple inner steps per outer iteration for stability

Deep Value Function Iteration (DVFI): Key Innovations

Lab 3B, Part 1 introduces three innovations beyond the basic actor-critic:

1. **Unified Objective (J):** T -step discounted sum plus terminal value:

$$J(y_0) = \sum_{t=0}^{T-1} \beta^t u(c_t) + \beta^T V(y_T)$$

2. **T -Step Lookahead:** Simulate the economy for $T = 30$ – 50 periods

- Gradients “see” long-term consequences of policy choices
- Reduces bias compared to one-step bootstrapping

3. **Soft Constraints** (penalty method):

- If $c > 0$: $u = \ln(c)$ (standard)
- If $c \leq 0$: $u = \text{Penalty} \times (c - \varepsilon)$ (pushes back to feasibility)
- Avoids hard clamping that kills gradients

DVFI: Separated Training Phases

- **Phase A — Policy Net** (maximize J):
 - Sample initial states y_0 uniformly in $[y_{\min}, y_{\max}]$
 - Simulate T periods using PolicyNet
 - Compute $J = \sum \beta^t u(c_t) + \beta^T V(y_T)$
 - Minimize $-J$ (gradient ascent on expected reward)
- **Phase B — Value Net** (fit V to J):
 - Target: J from Phase A, **detached** from policy gradients
 - Loss: $\|V(y_0) - J\|^2$ (standard regression)
- **Key detail:** `retain_graph=True` allows both phases to share the same computation graph; `detach()` in Phase B prevents policy updates through the value loss

Accuracy Benchmarks: Maximum Euler Residual

After training, evaluate the policy on a fine validation grid:

$$\log_{10}|EE(k_i, z_i)|, \quad i = 1, \dots, M$$

Standard benchmarks (Judd, 1998)


$\log_{10}|EE| < -3$: acceptable ($< 0.1\%$).

$\log_{10}|EE| < -5$: good ($< 10^{-3}\%$).

$\log_{10}|EE| < -7$: excellent.

Optimal growth (log utility, Cobb–Douglas):

- **Lab 3B Part 1 (DVFI, $T = 30$):** $\log_{10}|EE| \approx -3.5$ — slightly less precise but yields V too
- **Lab 3B Part 2 (Euler):** $\log_{10}|EE| \approx -5$ (64-64-1 PolicyNet, 5,000 epochs, AdamW)



The validation gate: dark = low residual (well trained); bright = boundary effects. Same shape Lab 3A/3B produces.

Comparison Against Lab 3A Ground Truth

- Analytical policy for log utility, $\delta = 1$:

$$c^*(y) = (1 - \alpha\beta) y, \quad k'^*(y) = \alpha\beta y$$

- Validate by plotting:
 - Trained policy $g_\theta(k)$ versus $k'^*(y)$ on the same axes
 - Pointwise error $|g_\theta(k) - k'^*(y)|$ as a function of k
 - Histogram of $\log_{10} |EE(k_i)|$ across a validation sample

- Sanity checks for Lab 3B (after 5,000 epochs):**

- Max policy error $< 10^{-3}$
- Savings rate $g(k)/y$ pinned at $\alpha\beta$ across the grid
- Loss curve monotone-decreasing on log scale, no oscillation

Next (T4): inject economic structure into the network — ICNN, monotone NNs — and the full Lab 3A (ground truth) / Lab 3B (NN) validation pipeline.