

Session 4.4

Optimal Taxation Without Commitment


Put a government into Deck 4.1's economy: a tax rule on the whole distribution, a Markov-perfect fixed point in rules, and the algorithm that finds it



Zhigang Feng

The pages marked PRELIMINARY are based on preliminary research by Zhigang Feng, Jiequn Han, Thomas J. Sargent, and Shenghao Zhu (September 2026). Do not circulate without the authors' permission.

Where Decks 4.1 and 4.2 Stopped



Deck 4.1, a household: $\mathbf{x}_{t+1}$ taken as given



this deck, a government: it sees $d\mathbf{x}_{t+1}$

- ▶ Deck 4.1: the histogram $\mathbf{x}$ and its matrix $\hat{\Phi}$, rebuilt at every date inside the graph. A household is atomistic: it takes $\mathbf{x}_{t+1}$ as given.
- ▶ Deck 4.2, *Argument 1*: a government that chooses its tax every period needs $d\Gamma'/d\tau$, and only a law of motion in the graph supplies it. Here τ is a function of $\mathbf{x}$, so households respond to it through $\mathbf{x}$.
- ▶ This deck puts that government into Deck 4.1's economy.
- ▶ **One object is the state, the law of motion and the gradient's path.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Decks 4.1 and 4.2.

The Question: a Tax Nobody Can Promise




- ▶ Under commitment the long-run capital tax is zero in the classic models: promise not to tax capital, and people save.
- ▶ Capital in place is an inelastic base. A government that chooses afresh each period is tempted to tax it; savers know it.
- ▶ Without commitment the problem is a game between today's government and its successors.
- ▶ The answer is a rule on the whole distribution, $\tau(\Gamma)$: not a number, not a function of a few moments.
- ▶ **Sought: the tax a government chooses when every successor will choose by the same rule.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Chamley (1986); Judd (1985); Klein, Krusell & Ríos-Rull, *RES* (2008).

Session 4.4

The Economy and Its Equilibrium

Deck 4.1's households, a government that taxes and cannot commit, and a fixed point in rules



Households

- ▶ As in Deck 4.1: ε redrawn every period (i.i.d.), $Q = 5$ Gauss–Hermite nodes (ε_q, p_q) , r net of depreciation. A household maximises

$$\sum_{t=0}^{\infty} \beta^t U(c_t, l_t, G_t), \quad U(c, l, G) = u\left(c - \psi_l \frac{l^{1+\theta_l}}{1+\theta_l}\right) + \gamma_g \log G, \quad u(y) = \frac{y^{1-\sigma}}{1-\sigma},$$


subject to

$$c_t + a_{t+1} = [1 + (1 - \tau_t^k) r_t] a_t + (1 - \tau_t^l) w_t \varepsilon_t l_t, \quad a_{t+1} \geq 0.$$

- ▶ τ^k taxes capital income, τ^l labour income; G is a public good, valued with weight γ_g .
- ▶ G enters additively, so it changes no household's choice: only the taxes do.
- ▶ **Hours still carry no wealth effect: the Frisch elasticity is $1/\theta_l = 3.125$.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. $\beta = 0.95$, $\sigma = 2$, $\theta_l = 0.32$, $\psi_l = 1.5$, $\sigma_\varepsilon = 0.2$.

Firms, a Balanced Budget, Market Clearing




no bonds: G is whatever the taxes raise this period

- ▶ Firm: $Y = K^\alpha L^{1-\alpha}$, $r = \alpha(K/L)^{\alpha-1} - \delta$, $w = (1 - \alpha)(K/L)^\alpha$.
- ▶ Government: picks (τ^k, τ^l) and balances its budget every period, $G = \tau^k rK + \tau^l wL$.
- ▶ Markets: $K = \sum_i x_i \bar{a}_i$, $L = \mathbb{E}[\varepsilon l(\varepsilon)]$, exactly; the histogram moves by $\mathbf{x}_{t+1} = \mathbf{x}_t \hat{\Phi}_t$.
- ▶ No bonds and no commitment technology: today's government cannot bind its successors.
- ▶ **G is a residual, not a choice: two instruments, one budget, every period.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. $\alpha = 0.34$, $\delta = 0.06$.

The Distribution, Approximated by a Histogram



- ▶ $N = 140$ bins on $[0, 16]$, midpoints $\bar{a}_1 < \dots < \bar{a}_N$; the state is the row vector $\mathbf{x} = (x_1, \dots, x_N)$, $x_i \geq 0$, $\sum_i x_i = 1$.
- ▶ Aggregates are sums: $K = \sum_i x_i \bar{a}_i$.
- ▶ ε is i.i.d.: $Q = 5$ Gauss-Hermite nodes (ε_q, p_q) integrate it, so $\mathbf{x}$ needs no ε dimension.
- ▶ **The government's state is these N numbers.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; as in Deck 4.1. The grid is fixed once, from the classical benchmark.

The Law of Motion $\mathbf{x}_{t+1} = \mathbf{x}_t \hat{\Phi}_t$, Built Entry by Entry

At date t , for every source bin i and node q

- 1 saving at the midpoint: $a'_{t,iq} = g_\theta(\varepsilon_q, \bar{a}_i; \mathbf{x}_t)$, kept inside $[\bar{a}_1, \bar{a}_N]$
- 2 bracket it, $\bar{a}_j \leq a'_{t,iq} \leq \bar{a}_{j+1}$ with $j = j_{t,iq}$, and weigh the upper neighbour: $\omega_{t,iq} = \frac{a'_{t,iq} - \bar{a}_j}{\bar{a}_{j+1} - \bar{a}_j}$
- 3 send $p_q(1 - \omega_{t,iq})$ of bin i 's mass to bin j and $p_q \omega_{t,iq}$ to bin $j + 1$ (Young's lottery)

$$\hat{\Phi}_{t,ij'} = \sum_{q=1}^Q p_q \left[(1 - \omega_{t,iq}) \mathbf{1}\{j' = j_{t,iq}\} + \omega_{t,iq} \mathbf{1}\{j' = j_{t,iq} + 1\} \right], \quad x_{t+1,j'} = \sum_{i=1}^N x_{t,i} \hat{\Phi}_{t,ij'}.$$

- Row i says where bin i 's households go: rows sum to one, so mass is conserved.
- Exact in the mean: $\sum_{j'} \hat{\Phi}_{t,ij'} \bar{a}_{j'} = \sum_q p_q a'_{t,iq}$, so K_{t+1} equals what households save.
- $\hat{\Phi}_t$ moves with θ and with $\mathbf{x}_t$, through a' : it is rebuilt at every date, and ω is linear in a' , so it can be differentiated.

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Young, JEDC (2010); Deck 4.1. $N \times Q$ network calls per date; no simulation, no sampling noise.

Household Conditions

- ▶ Saving: the Euler condition with the limit, equality whenever $a' > 0$,

$$u'(y_t) \geq \beta \mathbb{E}_{\varepsilon'} \left\{ [1 + (1 - \tau_{t+1}^k) r_{t+1}] u'(y_{t+1}) \right\}, \quad y = c - \psi_l \frac{l^{1+\theta_l}}{1 + \theta_l}.$$


- ▶ Hours, in closed form:

$$(1 - \tau^l) w \varepsilon = \psi_l l^{\theta_l} \implies l(\varepsilon) = \left(\frac{(1 - \tau^l) w \varepsilon}{\psi_l} \right)^{1/\theta_l}.$$

- ▶ *Tomorrow's* capital tax sits inside *today's* saving decision: expectations about successors move the economy now.
- ▶ Hours contain no α . Given τ^l and K , labour L , both prices and both tax bases are fixed, whatever the shape of $\mathbf{x}$.
- ▶ **Under GHH the marginal utility of consumption is $u'(y)$, so $U_c = u'(y)$ in every condition below.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Greenwood, Hercowitz & Huffman, *AER* (1988).

State, Rule, Timing



- ▶ State: $(\varepsilon, a; \mathbf{x})$ for a household; $\mathbf{x}$ alone for the government.
- ▶ A **Markov rule** $\Lambda(\mathbf{x}) = (\tau^k(\mathbf{x}), \tau^l(\mathbf{x}))$, with $G(\mathbf{x})$ from the budget: no calendar time, no history.
- ▶ Every government, at every date, uses the same rule, and that rule is the unknown.
- ▶ Households take a perceived rule $\tilde{\Lambda}$ and law of motion Φ as given, and best-respond.
- ▶ **The Markov restriction, not foresight, is what makes the rule time-consistent.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Klein, Krusell & Ríos-Rull (2008).

The Two Problems

(P-1) Household. Given prices, G , the perceived rule $\tilde{\Lambda}$ and $\mathbf{x}_+ = \mathbf{x}\hat{\Phi}$,

$$V(\varepsilon, a; \mathbf{x}) = \max_{0 \leq a'} \left\{ U(c, l, G) + \beta \sum_{q'} p_{q'} V(\varepsilon_{q'}, a'; \mathbf{x}_+) \right\}.$$

(P-2) Government (Deck 4.2's V^G). Given the successors' rule $\tilde{\Lambda}$ and its value $\tilde{W}$,


$$W(\mathbf{x}) = \max_{\tau^k, \tau^l} \left\{ \sum_{i,q} x_i p_q U(c_{iq}, l_q, G) + \beta \tilde{W}(\mathbf{x}_+) \right\},$$

subject to the budget, prices, the household's conditions at every node $(\bar{a}_i, \varepsilon_q)$, and $\mathbf{x}_+ = \mathbf{x}\hat{\Phi}$.

- ▶ Welfare is utilitarian: each node counts by its mass $x_i p_q$.
- ▶ **The asymmetry: the household takes $\mathbf{x}_+$ as given; the government is not atomistic and sees how its choice moves $\mathbf{x}_+$.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Klein, Krusell & Ríos-Rull (2008).

Markov-Perfect Equilibrium



- ▶ A **Markov-perfect equilibrium**: saving and value functions, prices, a rule Λ with value W , and perceptions, with
 - (a) households solving (P-1); (b) $\Lambda(\mathbf{x})$ solving (P-2) at every $\mathbf{x}$; (c) perceived = actual: $\tilde{\Lambda} = \Lambda$, $\tilde{W} = W$, Φ from households' saving; (d) markets clearing.
- ▶ No aggregate shocks: from *any* $\mathbf{x}_0$, taxes and aggregates must settle to constants.
- ▶ **A sharp test: paths from different starts must meet.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Schematic curves: direction and shape only.

Three Computational Challenges

Challenge	Deck 4.2's methods	Here
the distribution is a state	learned moments of a panel; a histogram; a history of shocks	Deck 4.1's histogram, and quadrature for ε
inequality constraints	Fischer–Burmeister: two inequalities become one equation	the same, at every node, priced by a penalty
strategic interaction	none of the three has a government that re-chooses	Deck 4.1's matrix, with the gradient passed through it

- ▶ Simulating Γ_+ works when nobody acts on the distribution; with a player who does, it breaks the forward and backward passes (Deck 4.2, *Argument 2*).
- ▶ **Deck 4.1 settled the first row. This deck takes the other two, then asks: is the answer right?**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Maliar et al. (2021); Azinovic et al. (2022); Han, Yang & E (2026).

Implementability: One Equation per Household State

- ▶ At every node $(\bar{a}_i, \varepsilon_q)$ the government must respect the Euler condition with the limit, as one equation:

$$\begin{aligned} \text{FB}(\vartheta_1, \vartheta_2) &= \vartheta_1 + \vartheta_2 - \sqrt{\vartheta_1^2 + \vartheta_2^2 + \eta^2} = 0, & \bar{E}_{iq} &= \text{FB}_{iq} / u'(y_{iq}), \\ \vartheta_1 &= u'(y) - \beta \sum_{q'} p_{q'} [1 + (1 - \bar{\tau}^k(\mathbf{x}_+)) r(\mathbf{x}_+)] u'(y_{+,q'}), & \vartheta_2 &= a'. \end{aligned}$$


- ▶ $\text{FB} = 0$ when both are non-negative and one is zero; $\eta = 10^{-4}$. $\bar{E}$ is a relative Euler error.
- ▶ a' is the saving *after* the clamp, as the budget and $\hat{\Phi}$ use it (Deck 4.1: pre-clamp $\tilde{a}$).
- ▶ Tomorrow's tax is the *perceived* rule $\bar{\tau}^k$, at the $\mathbf{x}_+$ today's saving produces.
- ▶ **The function is held fixed; the state it is evaluated at is not. There the strategic channel lives.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Maliar, Maliar & Winant (2021).

Session 4.4

The Ingredients and the Algorithm

Three networks, one penalised objective, a rollout through the matrix, and a frozen copy of the future




Three Trained Networks

Network	Input	Out	Head
saving g_θ	$(\varepsilon, a, \mathbf{x})$	1	Deck 4.1's change head, $\tilde{a} = a + D(2s - 1)$, $a' = \min\{\max\{\tilde{a}, 0\}, \bar{a}\}$
tax rule τ_ϕ	$\mathbf{x}$	2	a sigmoid times a cap: $\tau^k = \tau_{\max}^k \sigma(\cdot)$, $\tau^l = \tau_{\max}^l \sigma(\cdot)$
aggregate value $\bar{V}_\chi$	$\mathbf{x}$	1	linear: the government's value beyond the rollout
household critic V_ψ	$(\varepsilon, a, \mathbf{x})$	1	Deck 4.1's critic: loaded, not trained here

- ▶ Each: two hidden layers of 256 rectified-linear units. The tax rule reads the whole histogram, so it can respond to its *shape*.
- ▶ **The saving network reads $\mathbf{x}$, not the tax: in equilibrium the tax is itself a function of $\mathbf{x}$.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. $D = 1.3$, as in Deck 4.1.

The Government Chooses the Allocation Too




one backward pass,
two optimiser steps

- ▶ Deck 4.1 trained saving on the household's own Bellman equation. Here it learns only from the government's loss.
- ▶ The government picks the tax *and* the saving rule; $\bar{E} = 0$ makes it one household would choose.
- ▶ The **primal approach** of Ramsey taxation, with a Markov government in place of a committed one.
- ▶ **The rule is the households' only as far as the penalty drives $\bar{E}$ to zero.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Chari & Kehoe, *Handbook of Macroeconomics* (1999), ch. 26.


The Tax Rule's Head Decides What Can Be Learned



- ▶ $\tau^k(\mathbf{x}) = \tau_{\max}^k s_\tau$, $s_\tau = \sigma(\ell(\mathbf{x}))$, ℓ the last layer: its slope carries $\tau_{\max}^k s_\tau (1 - s_\tau)$.
- ▶ At the cap, $s_\tau \rightarrow 1$ and the gradient vanishes. A constant tax is a trivial fixed point: the run looks converged.
- ▶ Remedy one: a cap $\tau_{\max}$ wide enough for an interior optimum. A solver setting, not economics.
- ▶ Remedy two: a hinge barrier, zero inside, $B_\sigma = \sum [(0.02 - s_\tau)_+^2 + (s_\tau - 0.98)_+^2]$.
- ▶ **A floor under the gradient wherever τ lands, or the rule cannot learn to respond to $\mathbf{x}$.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.

Perceived Equals Actual: the Frozen Copy




today's choices are trained;
tomorrow's rule, saving and
value are read from the copy

- ▶ The frozen copy supplies $\bar{\tau}^k(\mathbf{x}_+)$ and $\bar{g}$ inside ϑ_1 , and $\bar{V}(\mathbf{x}_T)$ at the end of the rollout.
- ▶ Inside one step these perceptions do not move: they are targets, not parameters.
- ▶ Copying the live networks into the frozen ones imposes consistency: perceived = actual.
- ▶ It is Session 3's target network, now used to reach a fixed point in rules.
- ▶ **At convergence the rule the government expects of its successors is the rule it plays.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Session 3, *Bridge (3): the Household's Actor-Critic*.


The Loop in One Picture



- Steps 1–4 run four times per iteration, each on a fresh batch; steps 5–7 once, or at longer intervals.
- Steps 2–4 are one forward and one backward pass through $T = 30$ matrices.
- **No outer loop: the fixed point in rules is reached by training and synchronising.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.


Step 1a: Sampling Distributions From a Band



- ▶ Around an anchor $\mathbf{p}$, each bin may move by a share ν , $p_i(1-\nu) \leq x_i \leq p_i(1+\nu)$ (thin lines); the total stays one.
- ▶ **Chords**: free mass to the richest, or poorest, bins gives the highest, or lowest, K . Draw K uniformly on the line from $\mathbf{p}$ to that extreme.
- ▶ **Bin by bin**: each x_i uniform in its interval, projected back to mass one: every shape, little range in K .
- ▶ **Chords give reach in K , bin-by-bin draws give shape: the rule learns only where it is sampled.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Schematic: an illustrative anchor on 24 bins, $\nu = 0.35$.

Step 1b: The Band Follows, and the Policy's Own Orbit



- ▶ $\mathbf{x}^e, e = 1, \dots, E$: a batch of starting histograms, one rollout each. Half come from the band (Step 1a).
- ▶ Anchor $p_i \propto x_i^* e^{\kappa \bar{a}_i}$, moved every 10 iterations halfway to the classical K^* at current taxes: a taxing government moves both.
- ▶ Half from an **orbit**: every 20 iterations, $\mathbf{x} \leftarrow \mathbf{x} \hat{\Phi}(\mathbf{x}; \theta)$ for 60 steps, no gradient, plus the rule's own stationary $\mathbf{x}$.
- ▶ **The band covers off the path; the orbit is accurate on it.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.

Step 2: Roll Forward T Dates

From one sampled $\mathbf{x}_0 = \mathbf{x}^e$, for $t = 0, \dots, T - 1$


- 1 $K_t = \sum_i x_{t,i} \bar{a}_i$; taxes $(\tau_t^k, \tau_t^l) = \tau_\phi(\mathbf{x}_t)$; hours l_q , then L_t, r_t, w_t : no price loop
- 2 saving $a'_{t,iq} = g_\theta(\varepsilon_q, \bar{a}_i; \mathbf{x}_t)$; consumption $c_{t,iq}$ from the budget; $y_{t,iq}$
- 3 revenue $G_t = \sum_{i,q} x_{t,i} p_q [\tau_t^k r_t \bar{a}_i + \tau_t^l w_t \varepsilon_q l_q]$
- 4 flow welfare $\bar{U}_t = \sum_{i,q} x_{t,i} p_q [u(y_{t,iq}) + \gamma_g \log(G_t + 10^{-3})]$
- 5 build $\hat{\Phi}_t$ from the $a'_{t,iq}$ by Deck 4.1's lottery; $\mathbf{x}_{t+1} = \mathbf{x}_t \hat{\Phi}_t$, kept in the graph

$$W(\mathbf{x}^e; \theta, \phi) = \sum_{t=0}^{T-1} \beta^t \bar{U}_t + \beta^T \bar{V}(\mathbf{x}_T), \quad \bar{V} \text{ the frozen copy, a constant.}$$

- **Deck 4.1's rollout, with a tax at every date, the public good in welfare, and the matrix left in the graph.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. $T = 30$; utility floored far below any value reached, to guard the logarithm and the power.

Step 3: The Constraint at Every (i, q, t) , and Its Three Channels



each moves ϑ_1 ; all
three stay in the graph

- ▶ At each (t, i, q) and batch member:
 $\bar{E}_{t,iq} = \text{FB}(\vartheta_1, \vartheta_2)/u'(y_{t,iq})$.
- ▶ Inside ϑ_1 : $\mathbf{x}_+$ from the *live* saving rule; $\bar{\tau}$ and $\bar{g}$ from the frozen copy at $\mathbf{x}_+$.
- ▶ Deck 4.1 held a household's continuation fixed: it is its own marginal value (the envelope theorem).
- ▶ Here ϑ_1 is a *constraint*: no envelope theorem, and dropping a channel stalls the penalty.
- ▶ **The strategic channel, $\mathbf{x}_+ \rightarrow \bar{\tau}^k(\mathbf{x}_+)$, is the one no household sees.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.

Step 4: The Government's Loss

$$\mathcal{L}_\rho(\theta, \phi) = -W + \sum_{t,i,q} \beta^t x_{t,i} p_q \left[\mu_{t,iq} \bar{E}_{t,iq} + \frac{\rho}{2} \bar{E}_{t,iq}^2 \right] + B_{\text{feas}} + B_\sigma(\mathbf{x}^e) + \kappa_{\text{band}} B_{\text{band}}$$

- ▶ An augmented Lagrangian, multipliers μ . **As run, $\mu \equiv 0$: a pure quadratic penalty, weight ρ .**
- ▶ Weights $\beta^t x_{t,i} p_q$ price each residual by its mass; B_{feas} , same weights: $[(-c)_+ + (-y)_+] / 10^{-4}$.
- ▶ $B_{\text{band}} = \sum_t \beta^t \sum_i [(p_i(1 - \nu) - x_{t,i})_+ + (x_{t,i} - p_i(1 + \nu))_+]$, $\kappa_{\text{band}} = 20$: rollouts stay in the trained region.
- ▶ $B_\sigma = \sum_\tau [(0.02 - s_\tau)_+^2 + (s_\tau - 0.98)_+^2]$, $s_\tau = \tau / \tau_{\text{max}}$, at the sampled $\mathbf{x}^e$: zero inside, it keeps the sigmoid's slope alive, so the rule cannot freeze at its cap.
- ▶ No gradient through the weights, the $u'(y)$ dividing $\bar{E}$, or $\bar{V}(\mathbf{x}_T)$; one backward pass, one Adam step on θ and one on ϕ .
- ▶ **Every term but $-W$ is a solver device; the economics is $-W$ and $\bar{E} = 0$.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; Nocedal & Wright, *Numerical Optimization* (2006), ch. 17.

How a Network Responds to the Distribution

- ▶ The N bin masses enter the first layer like any other input. For the tax rule, two hidden ReLU layers:

$$\tau_{\phi}^k(\mathbf{x}) = \tau_{\max}^k \sigma(\ell_{\phi}(\mathbf{x})), \quad \ell_{\phi}(\mathbf{x}) = w_3^{\top} \text{ReLU}(W_2 \text{ReLU}(W_1 \mathbf{x}^{\top} + b_1) + b_2) + b_3.$$

- ▶ Its response to the mass in bin i , with D_1, D_2 the diagonal masks of active units:

$$\frac{\partial \tau_{\phi}^k}{\partial x_i} = \tau_{\max}^k \sigma'(\ell_{\phi}) [w_3^{\top} D_2 W_2 D_1 W_1]_i.$$

- ▶ Mass sums to one, so only moves count: δ from bin j to bin i changes τ^k by $\delta (\partial_i \tau^k - \partial_j \tau^k)$. A rise in K is one such direction; a mean-preserving spread is another.
- ▶ The saving rule $g_{\theta}(\varepsilon, \bar{a}_i; \mathbf{x})$ and the value $\bar{V}_{\chi}(\mathbf{x})$ read $\mathbf{x}$ the same way.
- ▶ $\partial \tau / \partial \mathbf{x}$ is learned, not imposed: it is how the government answers inequality, and it enters every saving update.

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Two hidden layers of 256 units; τ^l has its own output ℓ .

Updating the Tax Rule, One Date at a Time

$$\nabla_{\phi} \mathcal{L}_{\rho} = \sum_{t=0}^{T-1} \left[\frac{\partial \mathcal{L}_{\rho}}{\partial \tau_t^k} \nabla_{\phi} \tau_{\phi}^k(\mathbf{x}_t) + \frac{\partial \mathcal{L}_{\rho}}{\partial \tau_t^l} \nabla_{\phi} \tau_{\phi}^l(\mathbf{x}_t) \right] + \nabla_{\phi} B_{\sigma},$$

$$\frac{\partial \mathcal{L}_{\rho}}{\partial \tau_t^k} = -\beta^t r_t \left[\underbrace{\frac{\gamma_g K_t}{G_t + 10^{-3}}}_{\text{value of revenue}} - \underbrace{\sum_{i,q} x_{t,i} p_q u'(y_{t,iq}) \bar{a}_i}_{\text{cost to households}} \right] + \rho \sum_{i,q} \beta^t x_{t,i} p_q \bar{E}_{t,iq} \frac{\partial \bar{E}_{t,iq}}{\partial \tau_t^k}.$$

- ▶ g_{θ} does not read τ , so $\mathbf{x}_t$ does not depend on ϕ : each term sits at one date, at the $\mathbf{x}_t$ reached.
- ▶ The bracket is the flow trade-off, node by node; the penalty enters through today's $u'(y)$ in ϑ_1 .
- ▶ **Adam pushes τ_{ϕ} toward this condition at every $\mathbf{x}$ visited: the rule learns its response to the distribution only where the sampler and rollouts take it.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. r net; $\partial \bar{E} / \partial \tau_t^k$ through $u'(y_{t,iq})$, the divisor held fixed. τ^l : the same form, plus terms through hours, L , r , w .

Updating the Saving Rule: Through the Law of Motion

$$\mathbf{x}_{t+1} = \mathbf{x}_t \hat{\Phi}_t, \quad \hat{\Phi}_t = \hat{\Phi}(a'_t), \quad a'_{t,iq} = g_\theta(\varepsilon_q, \bar{a}_i; \mathbf{x}_t), \quad \frac{d\mathbf{x}_0}{d\theta} = 0.$$

$$\frac{d\mathbf{x}_{t+1}}{d\theta} = \underbrace{\frac{d\mathbf{x}_t}{d\theta} \hat{\Phi}_t}_{\text{displaced mass, carried on}} + \underbrace{\mathbf{x}_t \frac{\partial \hat{\Phi}_t}{\partial \theta}}_{\text{where today's mass lands}} + \underbrace{\mathbf{x}_t \frac{\partial \hat{\Phi}_t}{\partial \mathbf{x}_t} \frac{d\mathbf{x}_t}{d\theta}}_{\text{the saving rule reads } \mathbf{x}_t}$$

- Product rule on $\mathbf{x}_t \hat{\Phi}_t$, then chain rule: $\hat{\Phi}_t$ moves with θ directly and through $\mathbf{x}_t$.

Parameters	The gradient reaches them through
saving θ	consumption; the three terms, date after date; the taxes at each $\mathbf{x}_t$; the channels of ϑ_1
tax rule ϕ	the budget; G in welfare; $u'(y)$ in $\bar{E}$. Never $\mathbf{x}_+$: g_θ does not read τ
neither	$\bar{V}(\mathbf{x}_T)$; the frozen copy, a constant

- **Through θ , the gradient runs back along every $\hat{\Phi}_t$; through ϕ , only within the date.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. $\partial \hat{\Phi} / \partial a'$ is the lottery's slope, $\pm 1 / (\bar{a}_{j+1} - \bar{a}_j)$, zero where the clamp binds; the third term contracts over $\mathbf{x}_t$. Autograd runs the same chain in reverse.

Updating the Saving Rule: What the Distribution Sends Back

$$\lambda_t \equiv \frac{d\mathcal{L}_\rho}{d\mathbf{x}_t} = \frac{\partial \ell_t}{\partial \mathbf{x}_t} + \left(\lambda_{t+1} + \frac{\partial \ell_t}{\partial \mathbf{x}_{t+1}} \right) J_t, \quad J_t = \frac{\partial \mathbf{x}_{t+1}}{\partial \mathbf{x}_t}, \quad \lambda_T = 0,$$

$$\nabla_\theta \mathcal{L}_\rho = \sum_{t=0}^{T-1} \left[\frac{\partial \ell_t}{\partial \theta} + \left(\lambda_{t+1} + \frac{\partial \ell_t}{\partial \mathbf{x}_{t+1}} \right) \mathbf{x}_t \frac{\partial \hat{\Phi}_t}{\partial \theta} \right], \quad \ell_t : \text{date } t\text{'s terms of } \mathcal{L}_\rho.$$

In $\partial \ell_t / \partial \mathbf{x}_t$ who sits where, $\sum_q p_q U_{iq}$; prices via K_t ; **the tax rule's answer** $\partial \tau_\phi / \partial \mathbf{x}_t$; g_θ 's input
In $\partial \ell_t / \partial \mathbf{x}_{t+1}$ in $\bar{E}_t$: $\bar{\tau}^k(\mathbf{x}_{t+1})$, $r(\mathbf{x}_{t+1})$, $\bar{g}(\cdot; \mathbf{x}_{t+1})$: the strategic channel

- The tax rule's answer reaches θ twice: through $\tau_\phi(\mathbf{x}_t)$ along the path, and through $\bar{\tau}(\mathbf{x}_{t+1})$ in each Euler residual.
- **The saving rule is updated knowing how the tax will answer the distribution its own saving creates.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. The previous page run in reverse: one backward pass, $d\mathbf{x}_t/d\theta$ never formed. Penalty weights held fixed: "who sits where" enters through welfare only; $\bar{V}(\mathbf{x}_T)$ sends nothing back.

What Is New Here

	the state	who acts on $\mathbf{x}$	the gradient through $\mathbf{x}$'s motion
Krusell–Smith; Deck 4.2	mean, moments, histogram, shocks	nobody: households take it as given	one period, or a household panel
Markov taxes, one agent	K	a government	no distribution
HA taxes, commitment	a given $\mathbf{x}_0$	a government, once	a path of taxes, not a rule
this deck	the whole histogram	a government, every date	T dates, with $\partial\tau/\partial\mathbf{x}$ inside

- ▶ Without commitment the tax is a function $\tau(\mathbf{x})$; with commitment, a path chosen once.
- ▶ **To the authors' knowledge, the first time-consistent tax rule computed on the whole distribution as the state.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. One agent: Klein, Krusell & Ríos-Rull (2008). Commitment: Bhandari et al. (2021); Dyrda & Pedroni (2022).

Steps 5 and 7: Refit the Value Beyond the Horizon; Sync


- **Targets**, every 20 iterations, at fresh band draws $\mathbf{x}$. As run, flow welfare held forever:

$$\hat{V}(\mathbf{x}) = \frac{\bar{U}(\mathbf{x})}{1 - \beta}, \quad \bar{U}(\mathbf{x}) = \sum_{i,q} x_i p_q U_{iq}.$$

- **Fit**, every iteration: 20 Adam steps on $\text{mean}[\bar{V}_\chi(\mathbf{x}) - \hat{V}(\mathbf{x})]^2$, batches of 256, after a long prefit.
- **Sync**, every 10 iterations: $g_\theta, \tau_\phi, \bar{V}_\chi$ copied into $\bar{g}, \bar{\tau}, \bar{V}$. A hard copy, not a moving average.
- Beyond T the economy is valued as if it stayed where the rollout ended; at $\beta^T \approx 0.2$ that term is not negligible.
- **$\bar{V}$ supplies a level at the horizon, never a gradient; the rollout supplies the gradient.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.

Step 6: The Penalty Weight Follows the Residual



- ▶ Each iteration, read $\bar{E}^{\text{test}}$: the mass-weighted mean of $|\bar{E}|$ on a fixed test batch.
- ▶ Grow, $\rho \leftarrow \min(1.015\rho, 10^4)$, if the loss stalls ($|\Delta\mathcal{L}/\mathcal{L}| \leq 1\%$) while $\bar{E}^{\text{test}} \geq \text{tol}$ and not improving.
- ▶ Decay, $\rho \leftarrow 0.99\rho$, once 5 of the last 10 reads are below tol.
- ▶ Learning rates scaled by $1/\sqrt{1 + 0.1\rho}$.
- ▶ **Measure tol at the solution: below what is reachable, ρ grows without end.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Schematic curves: direction and shape only.

The Algorithm

Taxation without commitment, as a trained fixed point

- 1 warm-start g_θ from the no-tax solve; tax rule near zero; prefit $\bar{V}_\chi$; copy all to the frozen set
 - 2 four times: draw $\mathbf{x}^e$, half band, half orbit (step 1)
 - 3 roll forward T dates, $\hat{\Phi}_t$ in the graph; form W and every $\bar{E}_{t,iq}$ (steps 2–3)
 - 4 one backward pass of $\mathcal{L}_\rho$; one Adam step on θ , one on ϕ (step 4)
 - 5 refit $\bar{V}_\chi$; every 20 iterations new targets and orbit, every 10 a moved anchor (step 5)
 - 6 update ρ and the learning rates from the test-batch residual (step 6)
 - 7 every 10 iterations: frozen $\leftarrow$ live, “perceived = actual” (step 7)
 - 8 stop at the iteration budget, or early if the Euler error blows up; keep the best checkpoint
- **Each line is Deck 4.1’s machinery or Session 3’s target network. What is new is who owns the loss.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary.

What the Gradient Requires

Requirement	How it is met
a differentiable transition	lottery weights piecewise linear in a' (Deck 4.1); kinks of measure zero
no sampling noise	quadrature over ε , every asset node: a deterministic loss
a gradient in the tax rule	the barrier keeps the sigmoid off 0 and 1
a gradient at the limit	$\vartheta_2 = a'$ and the smoothed root keep constrained nodes live
memory for the graph	T dates of $N \times Q$ evaluations, $N \times N$ matrices, next-date copies in ϑ_1


- ▶ T is a budget. Beyond it, $\bar{V}$ adds a level and no gradient.
- ▶ Lab block B: derivatives through up to twenty lottery steps match finite differences to 2×10^{-10} .
- ▶ **Every row is a place where a gradient can vanish without an error message.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; the lab, block B.


Session 4.4

What the Rule Does Along a Path

Start the trained economy from different distributions and let the tax rule, saving and prices run




At the Same Capital, Inequality Alone Moves the Tax



- Five starts Γ_0 share one K_0 and differ only in inequality (a). At date 0, hours, prices and both tax bases are identical.
- The more unequal start is taxed less on impact (d); its capital overshoots further before returning (b, c).
- Every start converges to the same Γ^* and the same long-run tax.
- **A property of the trained rule: why it responds to shape is not yet settled.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Simulated with the trained Stage 2A networks.

At the Same Inequality, More Capital Means a Bigger Opening Levy



- ▶ The mirror experiment: the spread is held and K_0 moves (a). Now the tax base moves too, and the tax responds far more (d).
- ▶ The levy is transitional: capital and the interest rate return (b, c), and every start reaches the same long-run rate.
- ▶ A cap that never binds at the rest point can still censor the path: size it for the levy.
- ▶ **Without commitment: tax the capital inherited, then settle down.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary. Simulated with the trained Stage 2A networks.

Bridge: the Lab, and Session 5

Lab, Part I: the classical benchmarks of Decks 4.1–4.3.

Blocks D, E: Krusell–Smith by Deep Equilibrium Nets and in the sequence space.

Block F: the OLG economy, a market-clearing layer and a homotopy.

This deck: a government that differentiates through Γ .

- ▶ Cost: GPU-hours, not seconds. Stage 1B alone took about 2.9 GPU-hours, against 30 CPU-seconds for Stage 1A.
- ▶ It buys rules on the whole distribution, and the path from any start in the sampled region.
- ▶ Next in the research: aggregate shocks, $\hat{\Phi}$ conditioned on (z, z') ; then persistent ε .
- ▶ **The lab trains the Deck 4.2 and 4.3 networks in PyTorch, each graded by a gate fixed before the run.**

Source: Feng, Han, Sargent & Zhu (September 2026), preliminary; the Session 4 lab.